

# Attention: Motivation

Let's revisit the Encoder-Decoder architecture

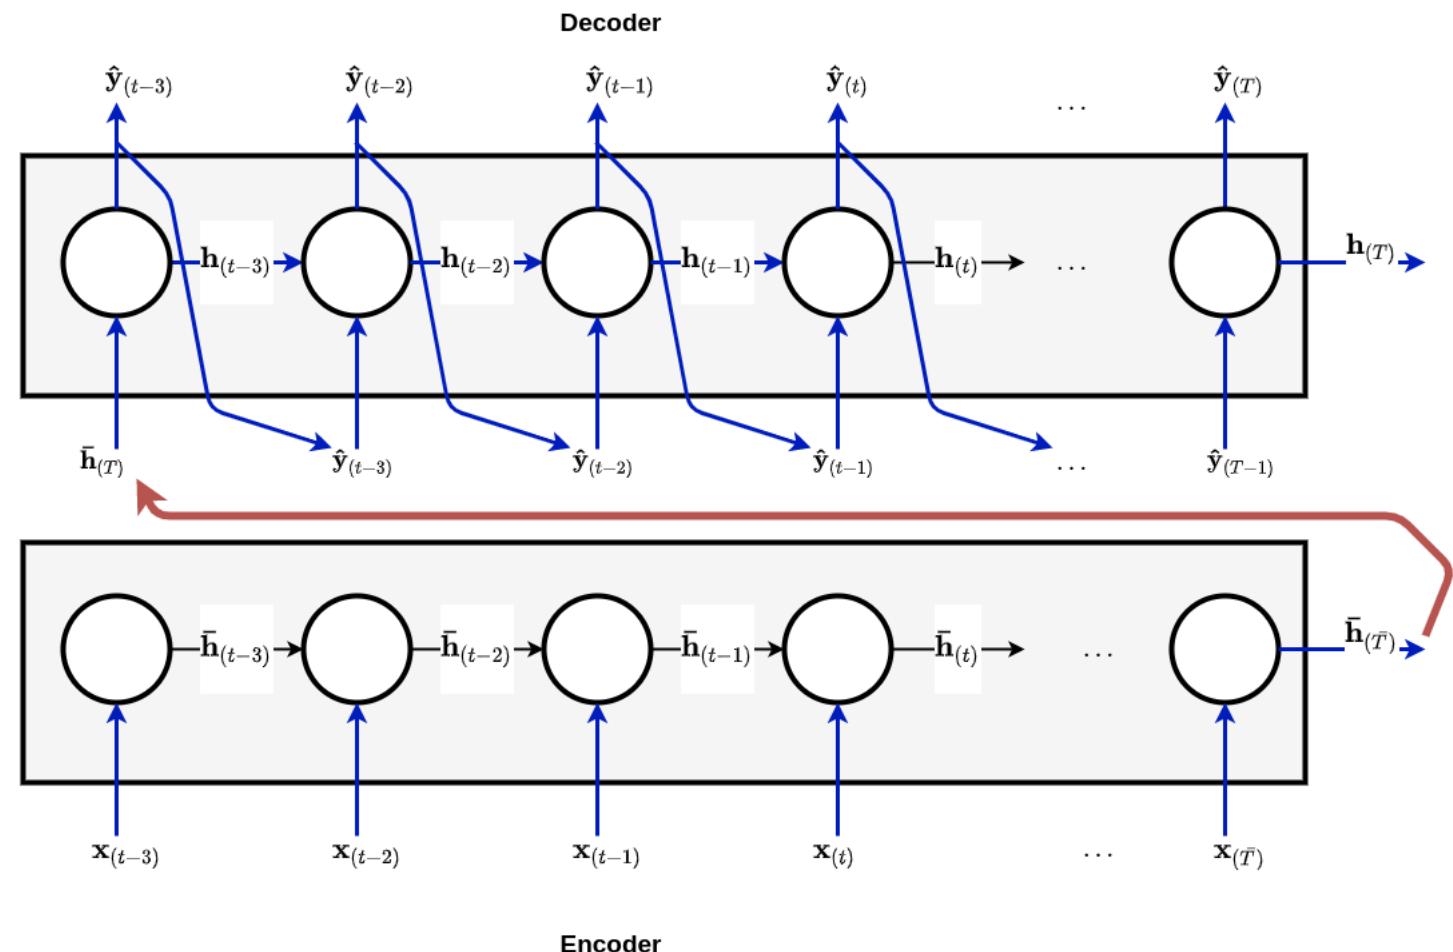
The Encoder

- Acts on input sequence  $[\mathbf{x}_{(1)} \dots \mathbf{x}_{(\bar{T})}]$
- Producing a sequence of latent states  $[\bar{\mathbf{h}}_{(1)}, \dots, \bar{\mathbf{h}}_{(\bar{T})}]$

## The Decoder

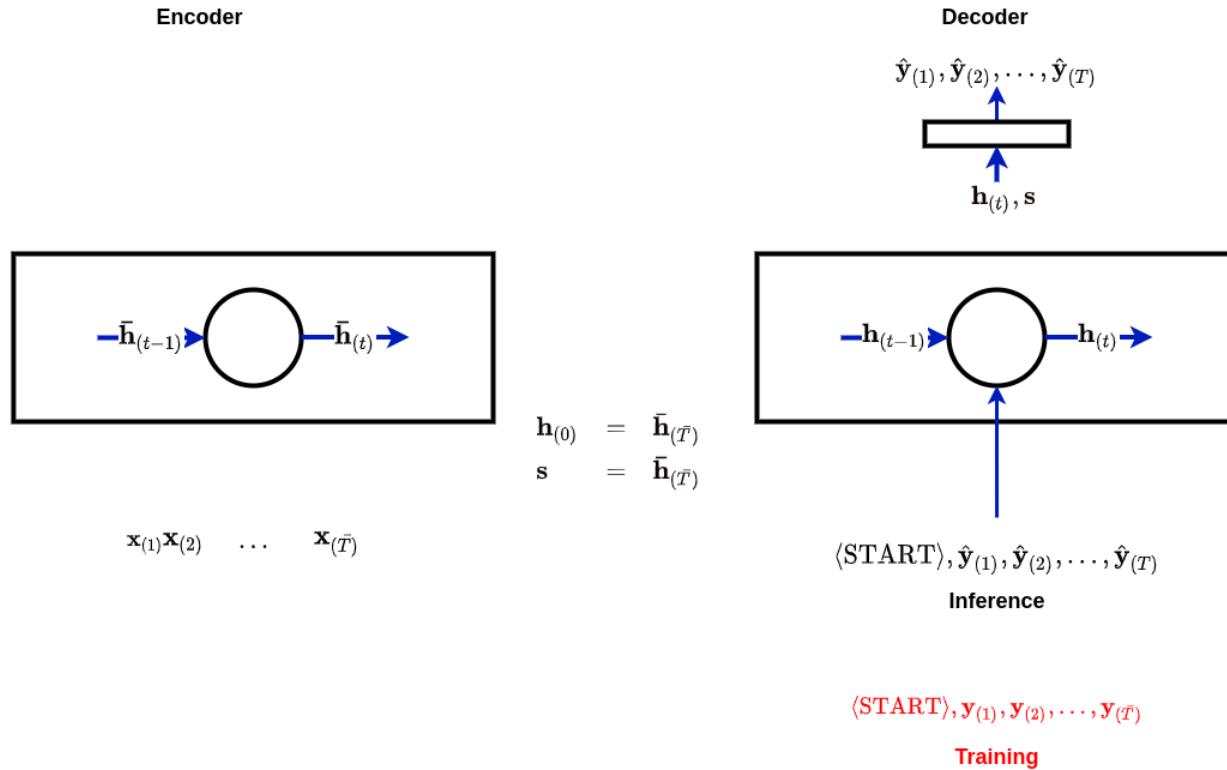
- Acts on the *final* Encoder latent state  $\bar{\mathbf{h}}_{(T)}$
- Producing a sequence of outputs  $[\hat{\mathbf{y}}_{(1)}, \dots, \hat{\mathbf{y}}_{(T)}]$
- Often feeding step  $t$  output  $\hat{\mathbf{y}}_{(t)}$  as Encoder input at step  $(t + 1)$

## RNN Encoder/Decoder



The following diagram is a condensed depiction of the process

Sequence to Sequence: training (teacher forcing) + inference: No attention



Recall that  $\bar{\mathbf{h}}_{(\bar{t})}$  is a fixed length encoding of the input prefix  $\mathbf{x}_{(1)}, \dots, \mathbf{x}_{(\bar{t})}$ .

So  $\bar{\mathbf{h}}_{(\bar{T})}$ , which initializes the Decoder, is a summary of entire input sequence  $\mathbf{x}$ .

This fact enables us to decouple the Encoder from the Decoder

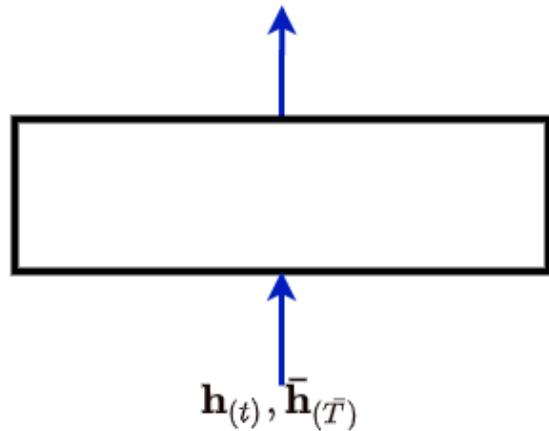
- The consumption of input  $\mathbf{x}$  and production of output  $\hat{\mathbf{y}}$  do not have to be synchronized
- Allowing for the possibility that  $T \neq \bar{T}$
- For example
  - There is no one to one mapping between languages (nor does ordering of words get preserved)

Let's focus on the part of the Decoder

- That transforms latent state (or short term memory)  $\mathbf{h}_{(t)}$  to output  $\hat{\mathbf{y}}_{(t)}$

**Decoder**

$$\hat{\mathbf{y}}_{(1)}, \hat{\mathbf{y}}_{(2)}, \dots, \hat{\mathbf{y}}_{(T)}$$



$$\bar{\mathbf{h}}_{(1)}, \bar{\mathbf{h}}_{(2)}, \dots, \bar{\mathbf{h}}_{(\bar{T})}$$

We can generalize this transformation as

$$\hat{\mathbf{y}}_{(t)} = D(\mathbf{h}_{(t)}; \mathbf{s})$$

In the vanilla RNN, this was governed by the equation

$$\hat{\mathbf{y}}_{(t)} = D(\mathbf{h}_{(t)}; \mathbf{s}) = \mathbf{W}_{hy} \mathbf{h}_{(t)} + \mathbf{b}_y$$

Additional parameter  $\mathbf{s}$

- Was unused in this example (our illustration used  $\bar{\mathbf{h}}_{(T)}$  as a place-holder)
- But may be used in other cases

This simple mapping of  $\mathbf{h}_{(t)}$  to  $\hat{\mathbf{y}}_{(t)}$  can be extremely burdensome

It is often the case that  $\hat{\mathbf{y}}_{(t)}$

- Depends mostly on a **specific element**  $\mathbf{x}_{(\bar{t})}$  of the input
- Or on a **specific prefix** of the input:  $\mathbf{x}_{(1)}, \dots, \mathbf{x}_{(\bar{t})}$

Consider the example of language translation

- When predicting word  $\hat{y}_{(t)}$  in the Target language
- Some "context" provided by the Source language may greatly influence the prediction
  - For example: gender/plurality of the subject

This context is usually much smaller than the entire sequence  $\mathbf{x}$  of length  $\bar{T}$ .

By not allowing  $D(\mathbf{h}_{(t)}; \mathbf{s})$  direct access to the required context, we force the Decoder

- To encode the context of the Source
- Along with the specific information of the Target
- Into  $\mathbf{h}_{(t)}$

This makes  $\mathbf{h}_{(t)}$  unnecessarily complex and perhaps difficult to learn well.

We will introduce a mechanism called *Attention* to alleviate this burden.

To give you a better feel for context, here are some examples

## Image captioning example

- Source: Image
- Target: Caption: "A woman is throwing a **frisbee** in a park."
- Attending over *pixels* **not** sequence

### Visual attention



---

A woman is throwing a **frisbee** in a park.

Attribution: <https://arxiv.org/pdf/1502.03044.pdf> (<https://arxiv.org/pdf/1502.03044.pdf>)

## Image captioning example

- Source: Image
- Target: Caption: "A giraffe standing in a forest with **trees** in the background."
- Attending over *pixels not sequence*

### Visual attention



---

A giraffe standing in a forest with **trees** in the background.

Attribution: <https://arxiv.org/pdf/1502.03044.pdf> (<https://arxiv.org/pdf/1502.03044.pdf>).

## Date normalization example

- Source: Dates in free-form: "Saturday 09 May 2018"
- Target: Dates in normalized form: "2018-05-09"

[link \(<https://github.com/datalogue/keras-attention#example-visualizations>\)](https://github.com/datalogue/keras-attention#example-visualizations)

# Attend to what's important

The solution to over-loading  $\mathbf{h}_{(t)}$  with Source context is conceptually straight forward.

In the Decoder expression  $D(\mathbf{h}_{(t)}; \mathbf{s})$ , let

$$\mathbf{s} = \mathbf{c}_{(t)}$$

where  $\mathbf{c}_{(t)}$  is a variable

- That supplies the appropriate context for output  $\hat{\mathbf{y}}_{(t)}$
- Conditional on  $\mathbf{h}_{(t)}$

Because  $\bar{\mathbf{h}}_{(\bar{t})}$

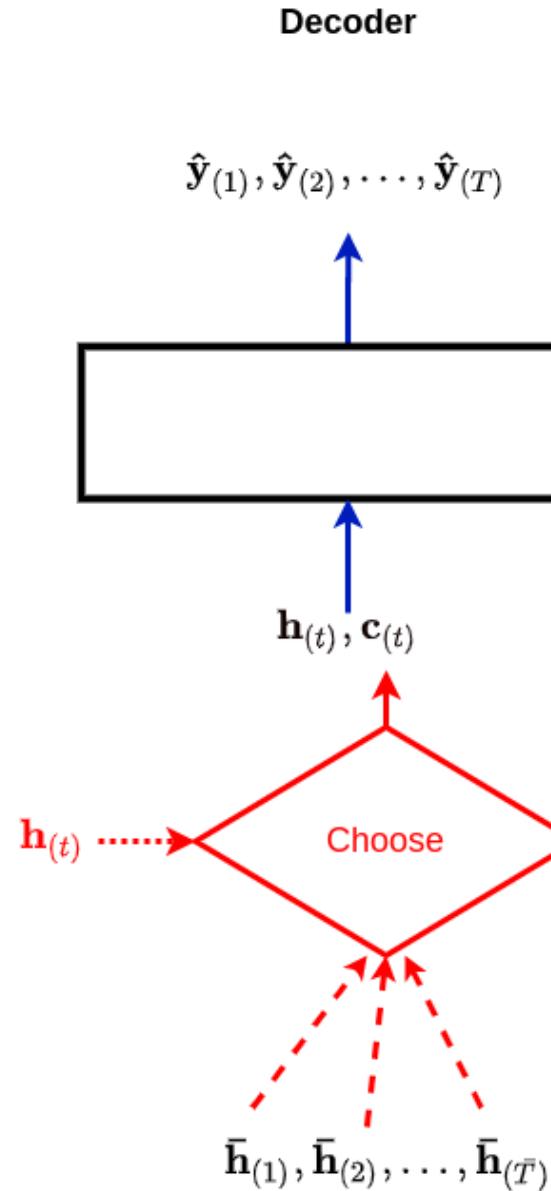
- Is a fixed length encoding of the input prefix  $\mathbf{x}_{(1)}, \dots, \mathbf{x}_{(\bar{t})}$
- It can be assigned to  $\mathbf{c}_{(t)}$  as the context for the prefix of  $\mathbf{x}$  of length  $\bar{t}$

$$\mathbf{c}_{(t)} \in \{\bar{\mathbf{h}}_{(1)}, \dots, \bar{\mathbf{h}}_{(\bar{T})}\}$$

We say

- The Decoder "attends to" (pays attention)  $\bar{\mathbf{h}}_{(\bar{t})}$
- When generating output  $\hat{\mathbf{y}}_{(t)}$

That is: it focuses its attention on a specific part of the input  $\mathbf{x}$





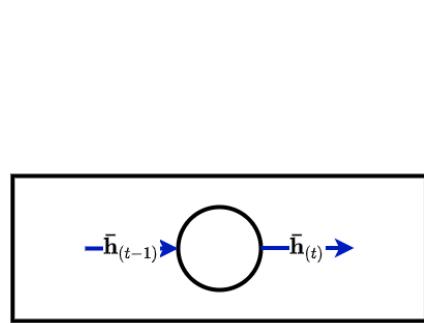
The dotted line from  $\mathbf{h}_{(t)}$  on the left of the Choose box

- Indicates that the Choice is conditional on Decoder state  $\mathbf{h}_{(t)}$

Here is a diagram summarizing the Attention mechanism

### Sequence to Sequence: attention

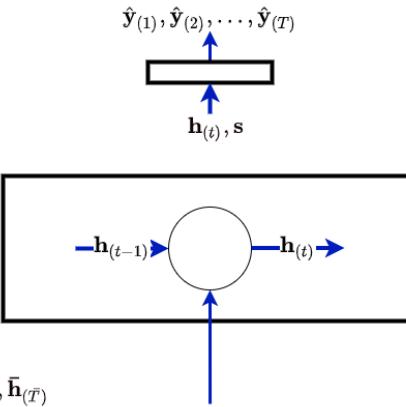
**Encoder**



$$\begin{aligned}\mathbf{h}_{(0)} &= \bar{\mathbf{h}}_{(\bar{T})} \\ \mathbf{s} &= \bar{\mathbf{h}}_{(1)}, \bar{\mathbf{h}}_{(2)}, \dots, \bar{\mathbf{h}}_{(\bar{T})}\end{aligned}$$

$\mathbf{x}_{(1)}, \mathbf{x}_{(2)}, \dots, \mathbf{x}_{(\bar{T})}$

**Decoder**



$\langle \text{START} \rangle, \hat{\mathbf{y}}_{(1)}, \hat{\mathbf{y}}_{(2)}, \dots, \hat{\mathbf{y}}_{(T)}$

**Inference**

$\langle \text{START} \rangle, \mathbf{y}_{(1)}, \mathbf{y}_{(2)}, \dots, \mathbf{y}_{(T)}$

**Training**

How is the choice of  $\mathbf{c}_{(t)}$  from the set  $\{\bar{\mathbf{h}}_{(1)}, \dots, \bar{\mathbf{h}}_{(\bar{T})}\}$  accomplished ?

The "Choose" box

- Is a Neural Network
- With its own weights
- That learn to make the best choice for the Target task !

In other words

- It is trained as part of the larger task

This is a common technique in Deep Learning that may, at first, appear magical

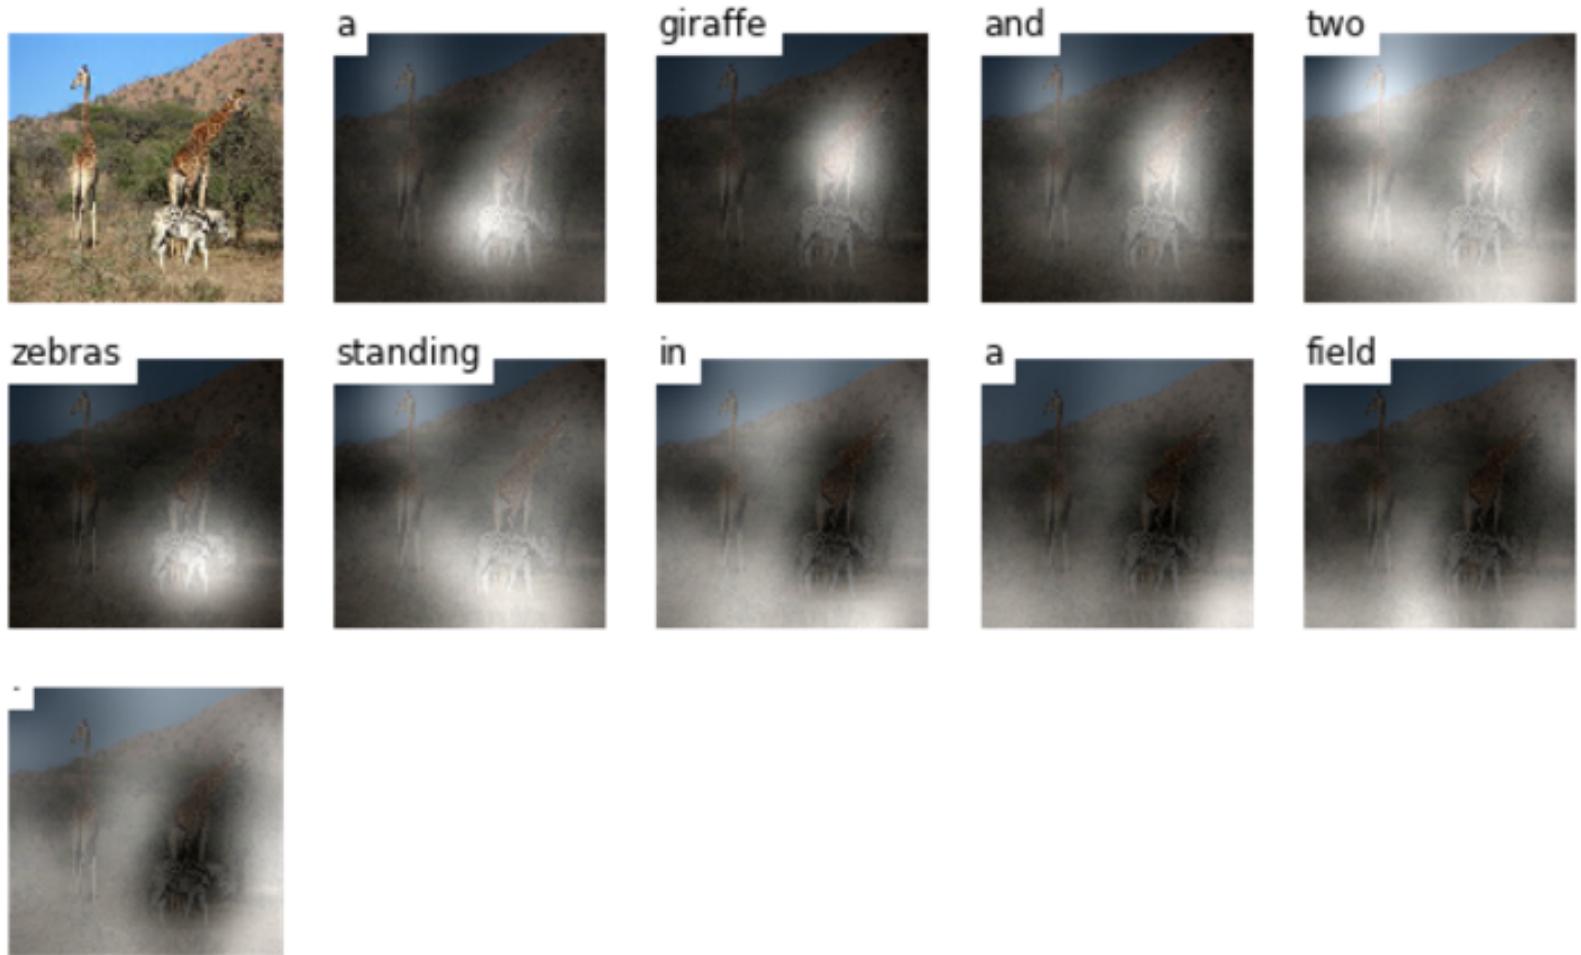
- Hypothesize the existence of a mechanism to solve your problem
- Train a Neural Network to conjure up the mechanism !

## **Just for fun: Attention in action**

Here are some examples of Sequence to Sequence problems using Attention.

## Visual Attention example

- Source: Image
- Target: Caption: "A giraffe and two zebras standing in a field."
- Attending over *pixels* **not** sequence



Attribution: <https://arxiv.org/abs/1502.03044> (<https://arxiv.org/abs/1502.03044>)

## Language Translation example

- Source: Spanish
  - Target: English
  - Colab notebook ! [Translation example](#)  
[\(https://colab.research.google.com/github/tensorflow/docs/blob/master/site/en/tut\)](https://colab.research.google.com/github/tensorflow/docs/blob/master/site/en/tut)
-

# Conclusion

We recognized that the Decoder function responsible for generating Decoder output  $\hat{\mathbf{y}}_{(t)}$

$$\hat{\mathbf{y}}_{(t)} = D(\mathbf{h}_{(t)}; \mathbf{s})$$

was quite rigid when it ignored argument  $\mathbf{s}$ .

This rigidity forced Decoder latent state  $\mathbf{h}_{(t)}$  to assume the additional responsibility of including Encoder context.

Attention was presented as a way to obtain Encoder context through argument  $\mathbf{s}$ .

In [2]: `print("Done")`

Done