

```
In [9]: # My standard magic ! You will see this in almost all my notebooks.
```

```
from IPython.core.interactiveshell import InteractiveShell  
InteractiveShell.ast_node_interactivity = "all"
```

```
# Reload all modules imported with %aimport
```

```
%load_ext autoreload
```

```
%autoreload 1
```

```
%matplotlib inline
```

## NumPy

[VandePlas Chapter 2 \(external/PythonDataScienceHandbook/notebooks/02.00-Introduction-to-NumPy.ipynb\)](#), [Geron notebook \(external/handson-ml/tools\\_numpy.ipynb\)](#).

## Python lists

Lists are *heterogeneous*: can contain elements of mixed type

```
In [1]: l = list( range(0,10) )  
        print(l)  
  
[0, 1, 2, 3, 4, 5, 6, 7, 8, 9]
```

```
In [2]: l[2] = "two"  
        print(l)  
  
[0, 1, 'two', 3, 4, 5, 6, 7, 8, 9]
```

Heterogeneity == *slow*

- Python interpreter has to constantly examine types

## NumPy ndarray

```
In [3]: import numpy as np
```

NumPy n-dimensional arrays ( `ndarray` ) are *homogenous*

- Can be faster because don't waste time examining type of each element
- Can be treated as vectors
- Vector arithmetic via compiled code = *fast*

```
In [4]: l = list( range(0,10))  
        l_plus_1 = [ e+1 for e in l]  
        print(l_plus_1)
```

```
[1, 2, 3, 4, 5, 6, 7, 8, 9, 10]
```

```
In [5]: l_np = np.array( np.arange(0,10))  
        print(l_np +1)
```

```
[ 1  2  3  4  5  6  7  8  9 10]
```

## Speed comparison

```
In [6]: list_len = 1000  
        l = list( range(0, list_len))  
        %timeit [ e+1 for e in l]
```

```
32.9 µs ± 132 ns per loop (mean ± std. dev. of 7 runs, 10000 loops each)
```

```
In [7]: l_np = np.array( np.arange(0, list_len) )  
        %timeit l_np +1
```

```
1.04 µs ± 6.71 ns per loop (mean ± std. dev. of 7 runs, 1000000 loops each)
```

**When dealing with large datasets, you need NumPy**

# Basics of NumPy arrays

[VanderPlas \(external/PythonDataScienceHandbook/notebooks/02.02-The-Basics-Of-NumPy-Arrays.ipynb\)](#)

[Vandeplas YouTube: Losing your loops - slides \(https://speakerdeck.com/jakevdp/losing-your-loops-fast-numerical-computing-with-numpy-pycon-2015?slide=14\)](#)

The most operation on ndarrays is indexing.

- ndarray indices are 0-based (i.e, first row/col is numbered 0, not 1)

```
In [10]: x = np.arange(0,6)

x

x[2]

M = np.arange(0,6).reshape(2,3)
M

M[1,1]
```

```
Out[10]: array([0, 1, 2, 3, 4, 5])
```

```
Out[10]: 2
```

```
Out[10]: array([[0, 1, 2],
               [3, 4, 5]])
```

```
Out[10]: 4
```

## Slicing

- Python (not just NumPy) upper bound of index is NOT inclusive

```
In [11]: print("x: ", x)
print("x tail: ", x[2:])
print("x head: ", x[:2])
```

```
x:  [0 1 2 3 4 5]
x tail:  [2 3 4 5]
x head:  [0 1]
```

## Strides

`x[start:stop:step]`

In [12]: `x[1:5:2]`

Out[12]: `array([1, 3])`

## Reshaping

In [12]: `grid = np.arange(1, 10).reshape((3, 3))`  
`print(grid)`

```
[[1 2 3]
 [4 5 6]
 [7 8 9]]
```

## Add dimensions

```
In [13]: x = np.arange(0,6)
print("x: ", x)
print("x shape: ", x.shape)

print("x re-shaped: ", x.reshape(1,-1))
print("x re-shaped shape: ", x.reshape(1,-1).shape)

print("x w/newaxis: ", x[ np.newaxis,:])
print("x w/newaxis sja[e: ", x[ np.newaxis,:].shape)
```

```
x: [0 1 2 3 4 5]
x shape: (6,)
x re-shaped: [[0 1 2 3 4 5]]
x re-shaped shape: (1, 6)
x w/newaxis: [[0 1 2 3 4 5]]
x w/newaxis sja[e: (1, 6)
```

## Concatentation, splitting

```
In [14]: x = np.array([1, 2, 3])
y = np.array([3, 2, 1])
x
y

np.concatenate([x, y])
```

```
Out[14]: array([1, 2, 3])
```

```
Out[14]: array([3, 2, 1])
```

```
Out[14]: array([1, 2, 3, 3, 2, 1])
```

You can concatenate multi-dimensional ndarrays:

```
In [15]: M1 = np.array([ [1, 2, 3],
                        [4, 5, 6]
                        ])

M2 = np.array([ [ 7,  8,  9 ],
                [10, 11, 12]
                ])

M1
M2

np.concatenate([ M1, M2 ])
```

```
Out[15]: array([[1, 2, 3],
                [4, 5, 6]])
```

```
Out[15]: array([[ 7,  8,  9],
                [10, 11, 12]])
```

```
Out[15]: array([[ 1,  2,  3],
                [ 4,  5,  6],
                [ 7,  8,  9],
                [10, 11, 12]])
```

You can also specify the dimension on which to concatenate



```
In [16]: M1  
         M2  
  
         np.concatenate([ M1, M2 ], axis=1)
```

```
Out[16]: array([[1, 2, 3],  
               [4, 5, 6]])
```

```
Out[16]: array([[ 7,  8,  9],  
               [10, 11, 12]])
```

```
Out[16]: array([[ 1,  2,  3,  7,  8,  9],  
               [ 4,  5,  6, 10, 11, 12]])
```

You can also use `vstack` (vertical stack) and `hstack` (horizontal stack)

```
In [17]: x = np.array([1, 2, 3])
          grid = np.array([[9, 8, 7],
                           [6, 5, 4]])

          y = np.array( [ [100],
                           [200]
                           ])

          x
          grid

          print("vstack:")
          # vertically stack the arrays
          np.vstack([x, grid])

          print("hstack:")
          y
          grid
          np.hstack( [y, grid])
```

```
Out[17]: array([1, 2, 3])
```

```
Out[17]: array([[9, 8, 7],
                [6, 5, 4]])
```

vstack:

```
Out[17]: array([[1, 2, 3],
                [9, 8, 7],
                [6, 5, 4]])
```

hstack:

```
Out[17]: array([[100],
                [200]])
```

```
Out[17]: array([[9, 8, 7],  
               [6, 5, 4]])
```

```
Out[17]: array([[100,  9,  8,  7],  
               [200,  6,  5,  4]])
```

## Ufuncs

[Vanderplasse \(external/PythonDataScienceHandbook/notebooks/02.03-Computation-on-arrays-ufuncs.ipynb#Introducing-UFuncs\)](https://vanderplasse.com/PythonDataScienceHandbook/notebooks/02.03-Computation-on-arrays-ufuncs.ipynb#Introducing-UFuncs).

### Math

- element-wise operations
- vectorized for speed
- operator overloading
  - $+$ ,  $-$ ,  $*$ ,  $/$
  - $<$ ,  $==$ ,  $>$
  - provides natural syntax
    - `l + 1`
    - ``np.add(l,1)`

```
In [18]: x = np.array( np.arange(0,10))
print("x: ", x)
print("+1: ", x + 1)
print("+1 verbose: ", np.add(x,1))
print("-1: ", x - 1)
```

```
x:  [0  1  2  3  4  5  6  7  8  9]
+1:  [ 1  2  3  4  5  6  7  8  9 10]
+1 verbose:  [ 1  2  3  4  5  6  7  8  9 10]
-1:  [-1  0  1  2  3  4  5  6  7  8]
```

## Aggregates

[Vanderplasse \(external/PythonDataScienceHandbook/notebooks/02.03-Computation-on-arrays-ufuncs.ipynb#Aggregates\)](https://vanderplasse.github.io/PythonDataScienceHandbook/notebooks/02.03-Computation-on-arrays-ufuncs.ipynb#Aggregates)

- Aggregation: taking a one-dimensional slice of an ndarray and reducing it to a scalar
  - also known as **reduce**

Best illustrated with an example

```
In [19]: x = np.arange(1, 6)
print("x: ", x)
print("x reduced by add: ", np.add.reduce(x))

# Less verbose synonym
print("x reduced by add, via sum", x.sum())
```

```
x: [1 2 3 4 5]
x reduced by add: 15
x reduced by add, via sum 15
```

## Aggregates on multi-dimensional ndarray: choose your dimension

```
In [20]: x = np.arange(1,7).reshape(2,3)
print("x: ", x)

print("x reduced on first dimension: ", x.sum(axis=0))
print("x reduced on second dimension: ", x.sum(axis=1))
```

```
x: [[1 2 3]
     [4 5 6]]
x reduced on first dimension: [5 7 9]
x reduced on second dimension: [ 6 15]
```

## Cumulative

Closely related to `reduce`: `accumulate`

- running operations, e.g, running sum

```
In [21]: print("x: ", x)
print("x running sum: ", np.add.accumulate(x)) # NOTE: not a method ON x; x is a parameter

# Less verbose synonym. n.b., WITHOUT an axis arg,, it will flatten x before summing
print("x running sum via cumsum: ", x.cumsum(axis=0))
```

```
x: [[1 2 3]
     [4 5 6]]
x running sum: [[1 2 3]
                [5 7 9]]
x running sum via cumsum: [[1 2 3]
                           [5 7 9]]
```

## Broadcasting

[Vanderplass \(external/PythonDataScienceHandbook/notebooks/02.05-Computation-on-arrays-broadcasting.ipynb\)](#)

You hopefully intuitively understand what NumPy does when a binary operator is applied to 2 identically-shaped arguments

```
In [22]: a = np.array([0, 1, 2])  
         b = np.array([5, 5, 5])  
         a + b
```

```
Out[22]: array([5, 6, 7])
```

But what happens if the two arguments have different shape ? Simplest case: one argument is dimension 0 or 1:

```
In [23]: print("a: ", a)  
         print("a + 1: ", a+1)
```

```
a:  [0 1 2]  
a + 1:  [1 2 3]
```

Next case: what if one argument is identical to the other EXCEPT is missing a dimension:

```
In [24]: M = np.arange(1,10).reshape(3,3)

print("a shape (", a.shape, "): ", a)
print("M shape (", M.shape, "):\n", M)
print("a + M shape(", (a+M).shape, "):\n", a + M)
```

```
a shape ( (3,) ): [0 1 2]
M shape ( (3, 3) ):
[[1 2 3]
 [4 5 6]
 [7 8 9]]
a + M shape( (3, 3) ):
[[ 1  3  5]
 [ 4  6  8]
 [ 7  9 11]]
```

NumPy took a one dimensional ndarray `a` and treated it like a 2-d ndarray by repeated it's rows

This is called **broadcasting**

Broadcasting follows some simple rules (quoted from [Vanderplass](#) ([external/PythonDataScienceHandbook/notebooks/02.05-Computation-on-arrays-broadcasting.ipynb](#))):