

```
In [1]: # My standard magic ! You will see this in almost all my notebooks.

from IPython.core.interactiveshell import InteractiveShell
InteractiveShell.ast_node_interactivity = "all"

# Reload all modules imported with %aimport
%load_ext autoreload
%autoreload 1

%matplotlib inline
```

Pandas

[VanderPlas Chapter 3 \(external/PythonDataScienceHandbook/notebooks/03.00-Introduction-to-Pandas.ipynb\)](#), [Geron notebook \(external/handson-ml/tools_pandas.ipynb\)](#)

```
In [2]: import numpy as np
import pandas as pd
```

Series

A Series is like a NumPy ndarray but with *symbolic indexing* (like a Dictionary)

```
In [3]: data = pd.Series([0.25, 0.5, 0.75, 1.0],  
                        index=['a', 'b', 'c', 'd'])  
data
```

```
Out[3]: a    0.25  
       b    0.50  
       c    0.75  
       d    1.00  
       dtype: float64
```

You select elements by using members of the index

```
In [4]: print("data at 'b': ", data["b"])
```

```
data at 'b': 0.5
```

A Series looks a little like a Dictionary, but with useful operations like NumPy ndarrays

```
In [5]: data.sum()
```

```
Out[5]: 2.5
```

You can access the index and the values directly. The values are a NumPy ndarray so you can easily integrate with NumPy.

The Index is it's own type.

```
In [6]: data.index  
  
data.values  
type(data.values)
```

```
Out[6]: Index(['a', 'b', 'c', 'd'], dtype='object')
```

```
Out[6]: array([0.25, 0.5 , 0.75, 1.  ])
```

```
Out[6]: numpy.ndarray
```

Symbolic indexing is super convenient ! No more "parallel" arrays of values and labels.

```
In [7]: data = pd.Series( np.arange(0,3), index=["2018-01-01", "2018-01-02", "2018-12-31"])  
data.index
```

```
Out[7]: Index(['2018-01-01', '2018-01-02', '2018-12-31'], dtype='object')
```

Pandas DataFrame

A DataFrame (note the capital "F") looks like a table or a 2-D ndarray but with *symbolic indexing* of rows and columns.

Unlike an ndarray, the columns can be different types.

Constructor

There are several ways to construct a DataFrame. Common methods:

- list of tuples, each tuple representing a row
- list of Dictionaries, each Dictionary representing a row (key/value pairs)
- A Dictionary
 - where keys are column names
 - values are the Series representing a column
 - the index of the Series becomes row names of the DataFrame

```
In [8]: prices      = pd.Series( { "FB": 150, "AAPL": 156, "AMZN": 1700, "NFLX": 340, "GOOG": 1100 } )
full_names = pd.Series( { "FB": "Facebook", "AAPL": "Apple", "AMZN": "Amazon",
                          "NFLX": "Netflix", "GOOG": "Google"
                          } )

stocks = pd.DataFrame({ "price": prices, "name": full_names })
stocks
```

Out[8]:

	price	name
FB	150	Facebook
AAPL	156	Apple
AMZN	1700	Amazon
NFLX	340	Netflix
GOOG	1100	Google

```
In [9]: ticker = "AMZN"
print( "Ticker {t}: price: {p}, full name: {fn}".format(t=ticker,
                                                         p=stocks.loc[ticker, "pr
ice"],
                                                         fn=stocks.loc[ticker, "n
ame"]
                                                         )
        )
```

Ticker AMZN: price: 1700, full name: Amazon

Notice how we indexed symbolically into the DataFrame: `stocks.loc[ticker, "price"]`

More on indexing.

Display

```
In [10]: stocks.head(2)

stocks.tail(3)
```

Out[10]:

	price	name
FB	150	Facebook
AAPL	156	Apple

Out[10]:

	price	name
AMZN	1700	Amazon
NFLX	340	Netflix
GOOG	1100	Google

Data Indexing and Selection

[VanderPlas \(external/PythonDataScienceHandbook/notebooks/03.02-Data-Indexing-and-Selection.ipynb#Data-Indexing-and-Selection\)](#)

There is more than one way to index, and it gets confusing !

Personally, I just use one-way: the `.loc` indexer

```
In [11]: print("Price column via .price attribute:\n", stocks.price)

# What would happen if "price" were in both the row index and the column index ?
print("\nPrice column via named column (implicitly the column):\n", stocks["price"])

print("\nPrice column via .loc indexer:\n", stocks.loc[:, "price"])
```

Price column via .price attribute:

FB	150
AAPL	156
AMZN	1700
NFLX	340
GOOG	1100

Name: price, dtype: int64

Price column via named column (implicitly the column):

FB	150
AAPL	156
AMZN	1700
NFLX	340
GOOG	1100

Name: price, dtype: int64

Price column via .loc indexer:

FB	150
AAPL	156
AMZN	1700
NFLX	340
GOOG	1100

Name: price, dtype: int64

Common "gotcha"

Note the difference in return type for the two slightly statements

```
In [12]: print("Arg. is an array with a single element")
stocks.loc[:, ["price"] ]
type(stocks.loc[:, ["price"] ] )

print("Arg. is a singleton")
stocks.loc[:, "price" ]
type(stocks.loc[:, "price" ])
```

Arg. is an array with a single element

Out[12]:

	price
FB	150
AAPL	156
AMZN	1700
NFLX	340
GOOG	1100

Out[12]: pandas.core.frame.DataFrame

Arg. is a singleton

Out[12]:

FB	150
AAPL	156
AMZN	1700
NFLX	340
GOOG	1100

Name: price, dtype: int64

Out[12]: pandas.core.series.Series

Index alignment

Indices are more than a convenient feature for indexing. They have semantic meaning in that they align two Series or DataFrames

```
In [13]: series1 = pd.Series( { "a": 1,    "b": 2,    "d": 4})  
        series2 = pd.Series( { "a": 10,    "c": 30,    "d": 44 })  
  
        series1  
        series2  
  
        series1 + series2  
        series2 + series1
```

```
Out[13]: a    1  
        b    2  
        d    4  
        dtype: int64
```

```
Out[13]: a    10  
        c    30  
        d    44  
        dtype: int64
```

```
Out[13]: a    11.0  
        b    NaN  
        c    NaN  
        d    48.0  
        dtype: float64
```

```
Out[13]: a    11.0  
        b    NaN  
        c    NaN  
        d    48.0  
        dtype: float64
```

Slices

You can use slices in Index !

Note that, unlike in Python, the upper bound is **inclusive**

Slices are particularly useful when you have DatetimeIndex (an index consisting of dates)

```
In [14]: stocks.loc["AMZN": "GOOG"]
```

Out[14]:

	price	name
AMZN	1700	Amazon
NFLX	340	Netflix
GOOG	1100	Google

Adding a column

```
In [15]: stocks["sector"] = pd.Series( { "FB": "Tech", "AAPL": "Tech", "AMZN": "ConsD",  
"NFLX": "ConsD" })  
stocks
```

Out[15]:

	price	name	sector
FB	150	Facebook	Tech
AAPL	156	Apple	Tech
AMZN	1700	Amazon	ConsD
NFLX	340	Netflix	ConsD
GOOG	1100	Google	NaN

Note the missing values.

Operations on Series and DataFrames

The usual suspects, plus vector operations like NumPy.

Note that the default `axis=0` , just like NumPy

```
In [16]: stocks.loc[:, "price"].mean()
```

```
Out[16]: 689.2
```

apply

There are lots of methods (hard to remember them all). Don't forget the old standby:
`apply`

The `apply` method to applies your own function, either column-wise or row-wise (depending on axis chosen)

```
In [17]: def my_func(s):  
        # If you don't understand what is passed (Series or DataFrame), or whether i  
        t it row or column, try this:  
        print("Type of s is {t}, shape is {sh}".format(t=type(s), sh=s.shape))  
  
        # s is a series  
        return s.mean()  
  
stocks.loc[:, ["price"] ].apply(my_func, axis=0)
```

Type of s is <class 'pandas.core.series.Series'>, shape is (5,)

```
Out[17]: price    689.2  
dtype: float64
```

Database (and Spreadsheet) like operations

Aggregation and grouping

[VanderPlas \(external/PythonDataScienceHandbook/notebooks/03.08-Aggregation-and-Grouping.ipynb\)](#)

The "split-apply-combine" paradigm ([VanderPlas \(external/PythonDataScienceHandbook/notebooks/03.08-Aggregation-and-Grouping.ipynb#GroupBy:-Split,-Apply,-Combine\)](#)) should be familiar to users of SQL, where the operation is called `group by`